Universität Trier, Fachbereich II
Professur Computerlinguistik

Layer Wise Gender Bias in Multilingual Foundation Models: A Five Architecture Cross Lingual Analysis

Machine Learning for Natural Language Understanding (MA2NLP2001)
Term Paper: Trends in Machine Learning

submitted by

Juan Jose Arango Serrano
Matr. 1861694
s2juaran@uni-trier.de

Hadil Athmouni
Matr. 1762515
s2haathm@uni-trier.de

Daniel Tesfai Kebede
Matr. 1716694
s4dakebe@uni-trier.de

Instructor: Raghvi Baloni

Abstract

Multilingual foundation models represent a central trend in modern machine learning, yet pretraining across massive corpora often encodes social biases. While previous studies have demonstrated that static embeddings distilled from contextualized models encode gender bias in English, Spanish presents a particular challenge as most professions are morphologically gender marked, confusing grammatical gender with social bias.

This study evaluates how static embeddings extracted from multilingual models assign gender bias to grammatically invariant Spanish professions (e.g., *dentista*, *taxista*) without explicit supervision. Comparative analysis is conducted across Arabic and Tigrina.

Using geometric test across multiple architectures, we demonstrate that while models successfully encode the grammatical neutrality of certain suffixes, statistical frequency and societal bias override morphological rules for fields historically dominated by men. Finally, our results demonstrate that the internal design of the model also has an influence on bias direction.

1 Introduction

Large language models trained across numerous languages represent a central trend in modern machine learning. Models such as mBERT and the XLM RoBERTa family create shared representational spaces for dozens of languages. This compression raises an important structural question, when languages with radically unequal training data coexist, how do these models distribute social associations across them? Bommasani et al. (2020) demonstrated that static embeddings distilled from English contextualized models encode social biases. However, English nouns generally lack grammatical gender. In Spanish, it is very common to find nouns marked by gender (e.g., *maestro/maestra*), which often confuses grammatical gender with social bias.

This research investigates how these models handle grammatically neutral Spanish professions compared to those with explicit gender markers. We also place Spanish within a broader comparative context alongside languages with lower training resources, such as Tigrinya and Arabic. By examining these models across their entire depth, we aim to understand how gender bias varies from the initial input to the final output. Previous studies establish that early layers process syntax while middle layers encode semantic relationships, which is exactly where we expect social associations to peak before a final representational collapse at the output layer.

Furthermore, we evaluate the robustness of our geometric bias measurements. We test whether a projection estimator sensitive to magnitude and a cosine estimator based entirely on angles agree on the underlying trends. We observed a Spearman correlation exceeding 0.96 across all valid models, confirming the reliability of our mathematical approach.

This research aims to fill the identified gap with a layer wise, cross linguistic and architecture comparative, with Spanish as the main focus, following the questions:

- **RQ1**: Do multilingual transformers encode measurable gender bias in Tigrinya, Arabic, and Spanish under a unified geometric protocol?
- **RQ2**: How does measured bias vary across transformer layers, and at which depths do associations peak?
- **RQ3**: Do two independent geometric estimators agree on bias trends across models?
- **RQ4**: Does attention architecture determine gender bias direction, independent of pretraining data?

2 Related Work

2.1 Bias in Static Word Embeddings

Early work established that social biases are encoded geometrically in static word embeddings. Bolukbasi et al. (2016) derived a gender direction from anchor sets, revealing systematic profession gender associations and proposing projection based debiasing. Caliskan et al. (2017) introduced WEAT as a statistical framework for implicit associations; Garg et al. (2018) showed embedding bias correlates with historical trends. Gonen and Goldberg (2019) demonstrated that removing a single bias direction does not eliminate deeper geometric bias, motivating continued geometric analysis.

2.2 Interpreting Contextual Representations via Reductions

Transformer models produce contextualised representations that vary across inputs and depth, complicating direct application of geometric methods. Bommasani et al. (2020) formalised reductions of contextual embeddings to stable word level vectors suitable for geometric analysis, while Rogers et al. (2020) surveys the unequal distribution of linguistic information across transformer layers, motivating layer wise analysis.

2.3 Multilingual Bias and Cross Lingual Considerations

As multilingual transformers became widely adopted, bias analysis extended beyond English. Zhao et al. (2019) and Zhao et al. (2020) showed that geometric bias methods transfer to contextual multilingual models and that bias patterns vary substantially across languages, with debiasing in one language not generalising to others. Guo and Caliskan (2021) further showed contextualised embeddings encode intersectional associations. These studies predominantly evaluate final layer representations. Lauscher et al. (2020) showed zero shot cross lingual transfer degrades substantially for typologically distant and low resource languages.

2.4 Architecture Driven Effects on Representations

The models evaluated here differ not only in scale but in fundamental attention design. He et al. (2021) introduced Disentangled Attention, which separates content and positional information into independent attention matrices, substantially changing how contextual relationships are encoded across layers. He et al. (2023) extended this design to multilingual pretraining via an ELECTRA-style training objective, producing the microsoft/mdeberta-v3-base checkpoint evaluated here. Liang et al. (2023) extended the Sentence Piece vocabulary from 250K tokens (XLM-R) to 1M tokens (XLM-V) to improve tokenisation coverage for low resource languages. These architectural and vocabulary decisions constitute the design space whose bias consequences this study directly compares.

2.5 How This Work Differs

We integrate geometric bias analysis Bolukbasi et al. (2016) with contextual-to-static reductions Bommasani et al. (2020) in a controlled layer wise, architecture comparative study across five multilingual models, revealing that Disentangled Attention produces a complete reversal of gender bias direction not previously documented in the multilingual bias literature.

3 Methodology

3.1 Models

We evaluate four pretrained multilingual transformer models without fine tuning. All checkpoints are loaded from HuggingFace Hub; no parameter updates are performed.

- **mBERT** (`bert-base-multilingual-cased`; Devlin et al., 2019): A 12-layer Transformer model (producing 13 hidden states including the embedding layer). It employs a WordPiece vocabulary of approximately 120K tokens covering 104 languages. It serves as our baseline for multilingual pre-training.
- **XLM-RoBERTa-base** (`xlm-roberta-base`; Conneau et al., 2020): Consists of 12 layers (13 hidden states) with a 250K-token SentencePiece vocabulary. It was trained on 2.5TB of CommonCrawl data, significantly scaling the training data compared to mBERT.
- **XLM-RoBERTa-large** (`xlm-roberta-large`; Conneau et al., 2020): A larger variant featuring 24 layers (25 hidden states) and a hidden dimension of 1024. It utilizes the same 250K SentencePiece vocabulary as the base model but provides higher capacity for complex linguistic patterns.
- **XLM-V-base** (`facebook/xlm-v-base`; Liang et al., 2023): This model maintains a 12-layer architecture (13 hidden states) but features an expanded SentencePiece vocabulary of 1M tokens. This design choice aims to solve the "vocabulary bottleneck" observed in low-resource language processing.
- **mDeBERTa-v3-base** (`microsoft/mdeberta-v3-base`; He et al., 2023): Comprising 12 layers (13 hidden states), this model utilizes *Disentangled Attention*. Unlike the standard Transformer architecture which sums content and position embeddings, DeBERTa represents them separately.

3.2 Languages and Data

This study focuses on Spanish (*es*). We utilize a sentence level corpus derived from **Spanish Wikimedia**, containing 100,000 sentences to provide a broad representation of natural language patterns.

Gender Anchors: To define the semantic gender direction, we use a set of anchor pairs (e.g., *él/ella*, *hombre/mujer*). These pairs provide the reference points for calculating the gender subspace in the embedding layers (Bolukbasi et al., 2016; Zhao et al., 2019).

Profession Lists: To separate properties from societal bias, we organize two distinct target lists of profession terms:

- **Group A (Invariant):** Grammatically neutral terms that remain identical regardless of the referent's gender (e.g., *dentista*, *comerciante*, *estudiante*, *periodista*).
- **Group B (Marked):** Grammatically gendered terms that possess distinct masculine and feminine forms (e.g., *ingeniero/ingeniera*, *abogado/abogada*). These serve as a baseline to establish the expected magnitude of the gender signal when it is explicitly encoded in the surface form.

3.3 Context Sampling

For each target word w, we retrieve up to n sentences containing w via substring matching. We use $n = 20$ contexts per profession term and $n = 12$ contexts per anchor term across all languages, to ensure stable centroid estimation within CPU memory constraints. Words with fewer than the minimum available contexts are excluded.

3.4 Contextual to Static Reduction

Transformer models produce contextualised token representations that vary across inputs. To enable geometric analysis, we adopt the reduction framework of Bommasani et al., 2020. For each occurrence of word w in context c_i, subword representations are mean-pooled. The result is averaged across n contexts. Formally, at layer ℓ:

$$s_\ell(w) = \frac{1}{n}\sum_{i=1}^{n} h_\ell(w, c_i), \tag{1}$$

where $h_\ell(w, c_i)$ is the mean pooled subword representation. This yields one stable static embedding per layer for each target word.

3.5 Gender Direction Construction

Following Bolukbasi et al., 2016, we define a gender direction using predefined anchor sets. Let M and F denote the male and female anchor term sets. For each layer ℓ, anchor representations are reduced following Eq. (1), and the centroids for the male and female anchor sets (A_M and A_F) are defined as:

$$c_{M,\ell} = \frac{1}{|M|}\sum_{w\in M} s_\ell(w), \quad c_{F,\ell} = \frac{1}{|F|}\sum_{w\in F} s_\ell(w), \tag{2}$$

and the layer specific gender direction is:

$$g_\ell = c_{M,\ell} - c_{F,\ell}. \tag{3}$$

This vector encodes the principal geometric contrast between male and female anchors at layer ℓ (Bolukbasi et al., 2016).

3.6 Bias Measurement and Agreement

For each profession p and layer ℓ we compute two geometric estimators:

(a) **Gender Direction Projection (`proj`)**: The scalar projection of $s_\ell(p)$ onto the gender direction:

$$ProjBias_\ell(p) = \frac{s_\ell(p)\cdot g_\ell}{\|g_\ell\|}. \tag{4}$$

Positive values indicate alignment toward the male centroid; negative values indicate alignment toward the female centroid. The scalar projection preserves vector magnitude and is sensitive to the strength of the gender association (Bolukbasi et al., 2016).

(b) **Centroid Cosine Difference (`cosdiff`)**. As a complementary magnitude normalised estimator:

$$\text{CosBias}_\ell(p) = \cos(\mathbf{s}_\ell(p), \mathbf{c}_{M,\ell}) - \cos(\mathbf{s}_\ell(p), \mathbf{c}_{F,\ell}), \tag{5}$$

where cosine similarity is defined as:

$$\cos(\mathbf{a}, \mathbf{b}) = \frac{\mathbf{a}\cdot\mathbf{b}}{\|\mathbf{a}\|\|\mathbf{b}\|}. \tag{6}$$

This estimator directly compares relative angular proximity to male and female centroids and is invariant to

vector scale.

3.7 Estimator Agreement

To assess measurement robustness, we compute Spearman rank correlation (Spearman, 1904) between the two estimators across profession terms at each layer:

$$\rho_\ell = \text{Spearman}(\{ProjBias_\ell(p)\}_p, \{CosBias_\ell(p)\}_p). \tag{7}$$

$\rho\ell = 1$ indicates perfect ordinal agreement across professions.

3.8 Layer Wise Aggregation

For visualisation we compute the mean projection score over profession terms at each layer ℓ, producing a layer wise bias trajectory:

$$\overline{\text{Proj}}_\ell = \frac{1}{|P|} \sum_{p \in P} \text{ProjBias}_\ell(p), \tag{8}$$

where P is the set of professions. This trajectory summarises how the geometric gender signal evolves across model depth.

4 Results

4.1 Estimator Agreement (RQ3)

Before interpreting bias scores, we verify measurement robustness. Table 1 reports the Spearman rank correlation ρ between the projection and cosine-difference estimators for each model across all layers. Every model achieves a minimum $\rho \geq 0.97$ and a mean above 0.99, showing a near perfect agreement. The one notable exception is XLM-V at layer 12, where ρ drops to 0.70, consistent with the representational collapse commonly observed at the final layer.

Table 1: Estimator agreement: Spearman ρ between projection and cosine-difference scores across layers (Spanish).

Model	**Min ρ**	**Max ρ**	**Mean ρ**
XLM-R Base	0.991	1.000	0.997
XLM-R Large	0.986	1.000	0.997
XLM-V Base	0.700	1.000	0.976
mDeBERTa-v3	0.967	0.999	0.990
mBERT	0.979	0.998	0.990

4.2 Group A: Invariant Professions (RQ1)

Table 2 reports projection scores for the ten invariant professions. We report measurements at each model's peak layer (e.g., Layer 9 for XLM-R Base) where bias magnitude is highest. As shown in Figure 3, computing a mean across all layers dilutes the signal because early layers focus on syntax and final layers prepare logits for output. The full layer by layer trajectories are provided in subsequent heatmaps.

Some terms show neutral morphology: *dentista* (+0.01 in XLM-R Base) and *cantante* (−0.21 in XLM-R Large) remain near zero. However, other invariant terms exhibit strong male bias. For instance, *comerciante* scores +2.69 in XLM-R Base, while mDeBERTa classifies it neutrally (−0.03). *Piloto* shows consistently high male bias across all evaluated models.

Table 2: Projection scores for Group A (invariant professions) at each model's peak layer. Positive values indicate male bias.

Profession	**XLM-R B (L9)**	**XLM-R L (L16)**	**XLM-V (L7)**	**mDeBERTa (L5)**	**mBERT (L11)**
policía	+2.09	+0.85	−0.77	−0.22	+1.12
artista	+3.08	+2.14	+1.49	+0.64	+5.04
comerciante	+2.69	+1.41	+2.36	−0.05	+3.71
periodista	+2.56	+1.03	+0.81	+0.26	+3.13
piloto	+3.33	+1.78	+2.03	+1.22	+4.72
cantante	+2.15	−0.21	+0.01	+0.31	+0.36
juez	+2.56	+1.96	+1.19	+0.76	+3.32
dentista	+2.08	+1.41	+1.02	+0.19	+3.00
estudiante	+2.40	+1.74	+1.35	+0.76	+4.43

4.3 Group B: Gendered Profession Pairs (RQ1)

Table 3 presents peak layer projection scores for the masculine and feminine forms of some of the gendered professions, alongside their magnitude asymmetry ($\Delta = |proj_M| - |proj_F|$). Figure 1 visualises this geometry for XLM-R Base.

While masculine forms align positively and feminine forms align negatively, the absolute magnitude of masculine scores is consistently higher. For example, in XLM-R Base at layer 9, *profesor* scores +3.45 while *profesora* scores +0.08 ($\Delta = 3.37$). This asymmetry indicates that masculine forms serve as the default geometric anchor for the profession concept, effectively treating feminine forms as semantic deviations. Figure 2 further demonstrates that this massive disparity between masculine and feminine projections is not an isolated artifact of the peak layer, but a persistent mathematical reality across the entire depth of the architecture.

Table 3: Peak layer projection scores for Group B gendered pairs (XLM-R Base, Layer 9). $\Delta = |proj_M| - |proj_F|$. Positive Δ indicates masculine dominance.

Profession	**Masculine**	**Feminine**	**Δ**
abogado/a	+3.48	+0.68	+2.80
profesor/a	+3.45	+0.08	+3.37
mecánico/a	+2.94	+1.15	+1.80
cocinero/a	+2.90	+0.56	+2.34
minero/a	+2.66	+0.90	+1.76
arquitecto/a	+3.06	+1.52	+1.54
ingeniero/a	+3.15	+0.88	+2.27

4.4 Layer Wise Bias Trajectory (RQ2)

Figure 3 shows the mean projection score across representations at each transformer layer. XLM-R Base, XLM-R Large, and XLM-V share a rise peak decay trajectory: bias is weak initially, peaks in the middle layers (7 to 16), and decreases toward output generation. This matches observations that middle layers encode semantic relationships.

In contrast, mDeBERTa-v3-base never develops the male biased signal seen in RoBERTa models but instead, its projection remains near zero or slightly negative across all layers. This architectural shift is analysed in

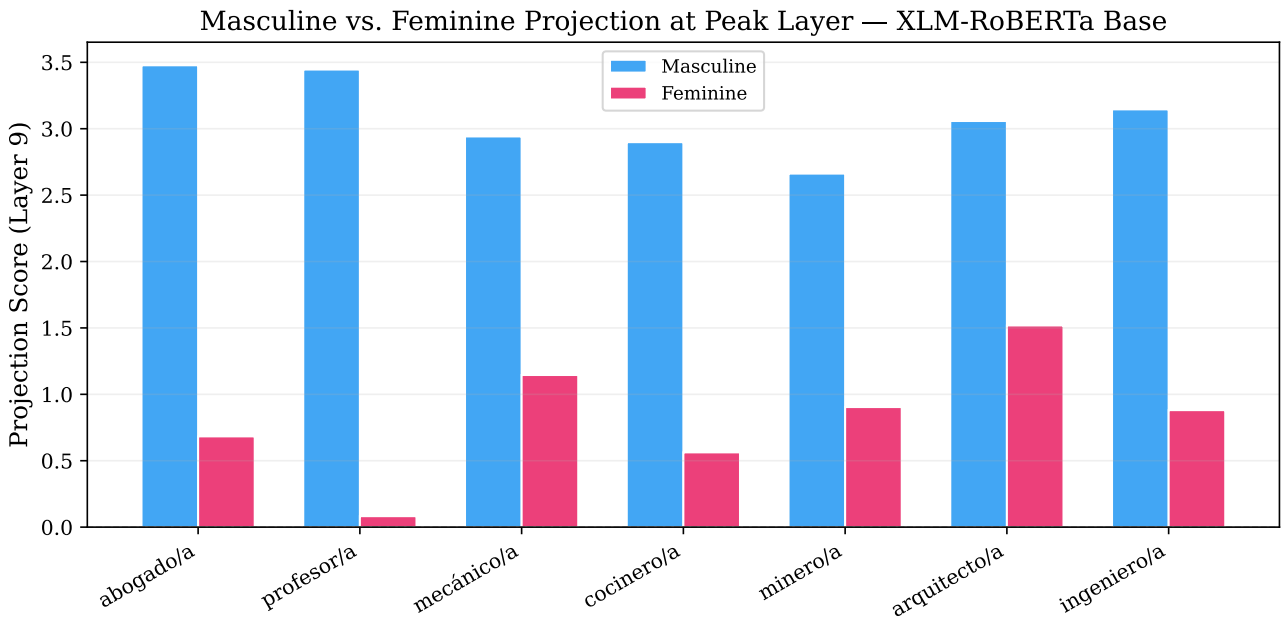

Figure 1: Masculine vs. feminine projection scores at peak layer 9 for XLM-RoBERTa Base.

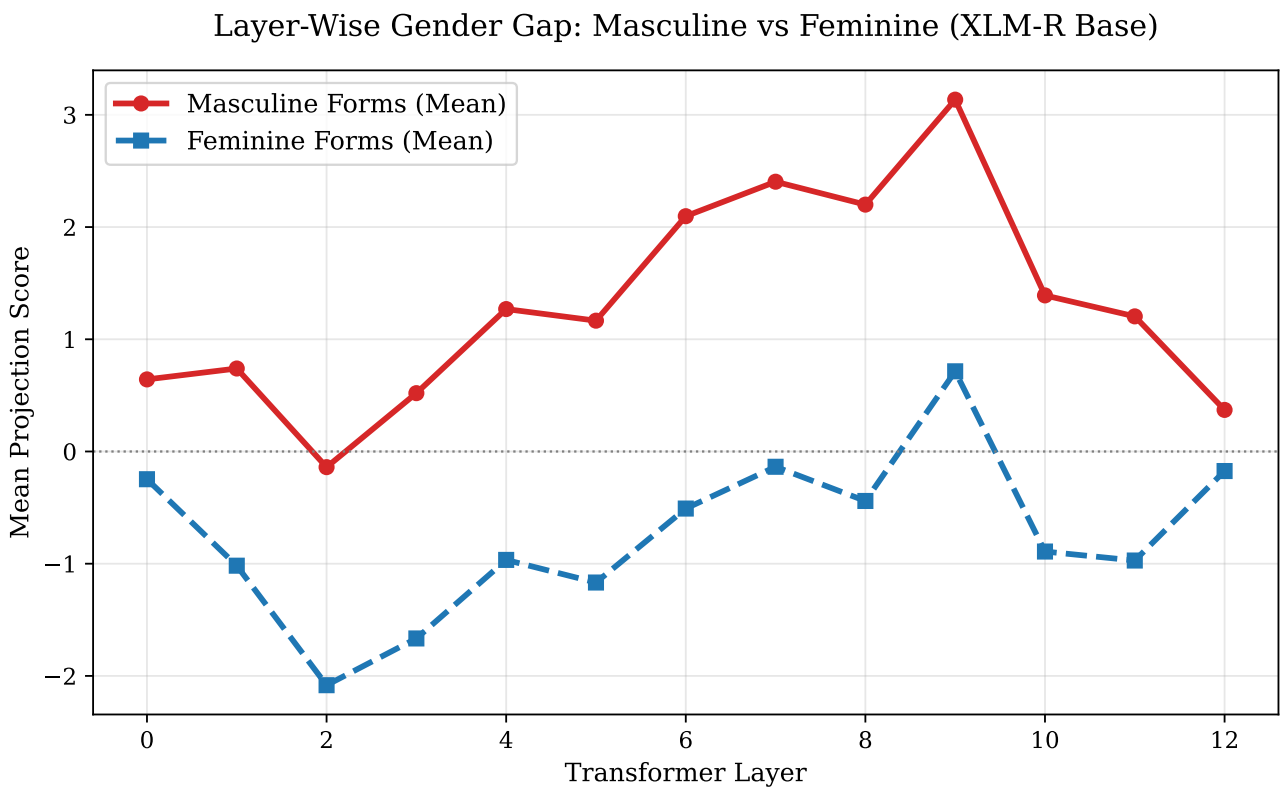

Figure 2: Layer wise mean projections for masculine vs. feminine forms in XLM-RoBERTa Base. The masculine forms persistently dominate the semantic space across all network depths.

Section 5.

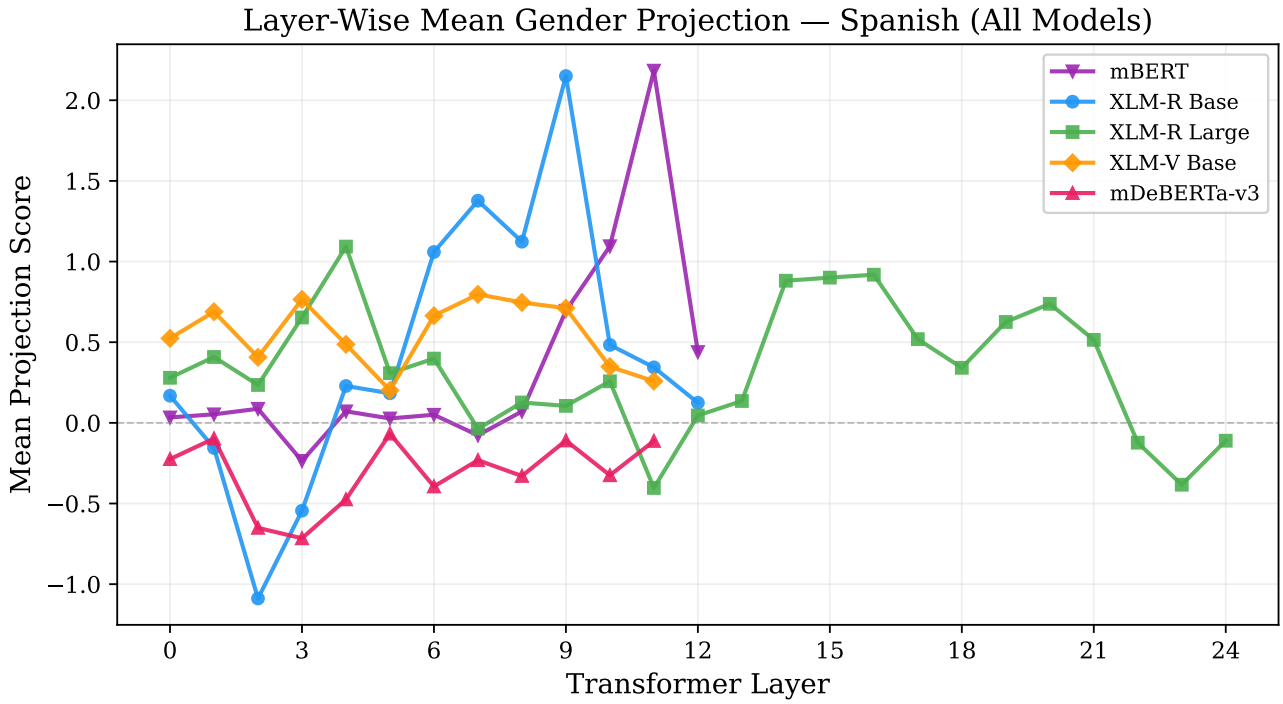

Figure 3: Layer wise mean gender projection across all professions for four Spanish models. The three RoBERTa family models show a rise, peak and decay pattern while mDeBERTa-v3 stays near or below zero.

4.5 Architecture-Driven Direction Reversal (RQ4)

Figure 4 contrasts per profession heatmaps for XLM-R Base and mDeBERTa-v3. While XLM-R Base displays positive male values in middle layers, mDeBERTa projects predominantly negative or near zero values.

Since both models use comparable training corpora, this reversal likely results from mDeBERTa's *Disentangled Attention* mechanism. By separating content and positional matrices, the model mitigates the accumulation of gender associations from grammatical context into the core representations.

The full per-profession heatmap for XLM-R Base (Figure 5) illustrates the layered structure: Group A terms cluster near the boundary between red and blue, while Group B masculine terms are deep red and feminine terms are deep blue.

4.6 Cross Lingual Patterns (RQ1)

To contextualise these results, Figure 6 maps the layer wise projection for XLM-RoBERTa Base across Spanish, Arabic, and Tigrinya, enabling a controlled comparison. Table 4 provides key metrics.

Despite large disparities in pretraining corpus size and morphological structure, the bias trajectory is very consistent. For all three languages, XLM-R Base produces a smooth rise that peaks precisely at Layer 9, with magnitudes bounded between +2.15 and +2.99. This shows that the architecture embeds gender bias in the same way, regardless of the language's resource level.

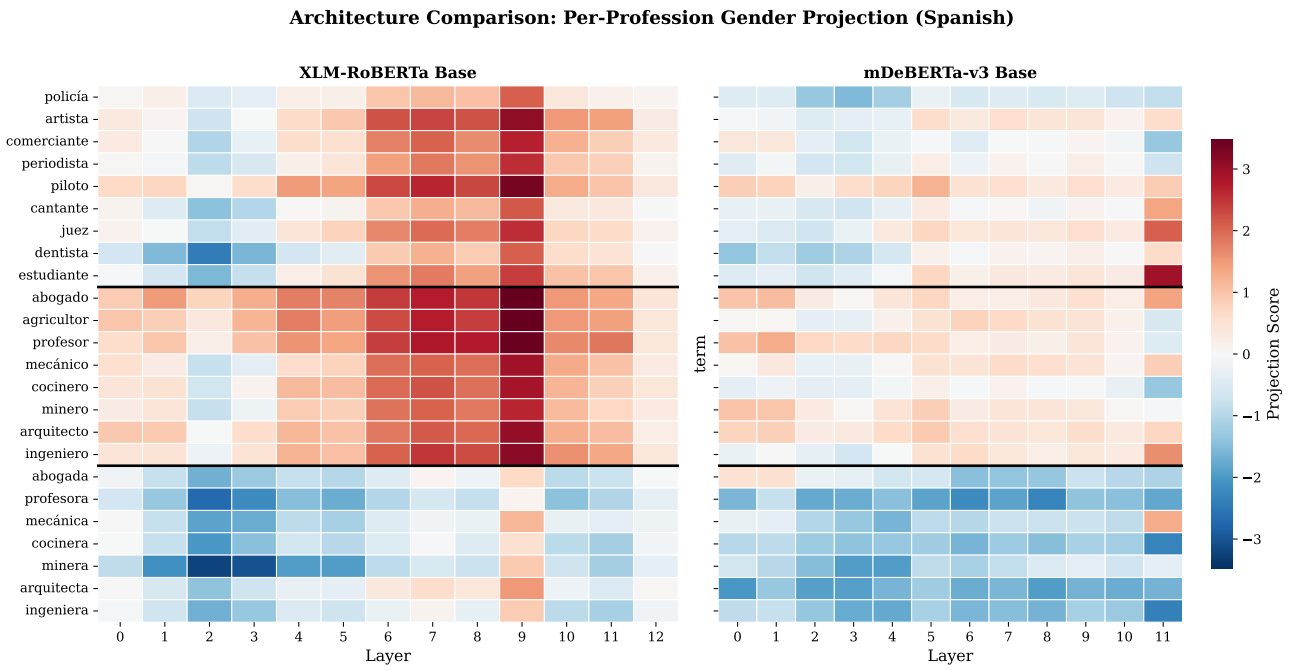

Figure 4: Per profession projection heatmaps: XLM-RoBERTa Base (left) vs. mDeBERTa-v3 (right). Disentangled Attention produces a reversal in bias direction across all professions. Layer 12 is excluded for mDeBERTa.

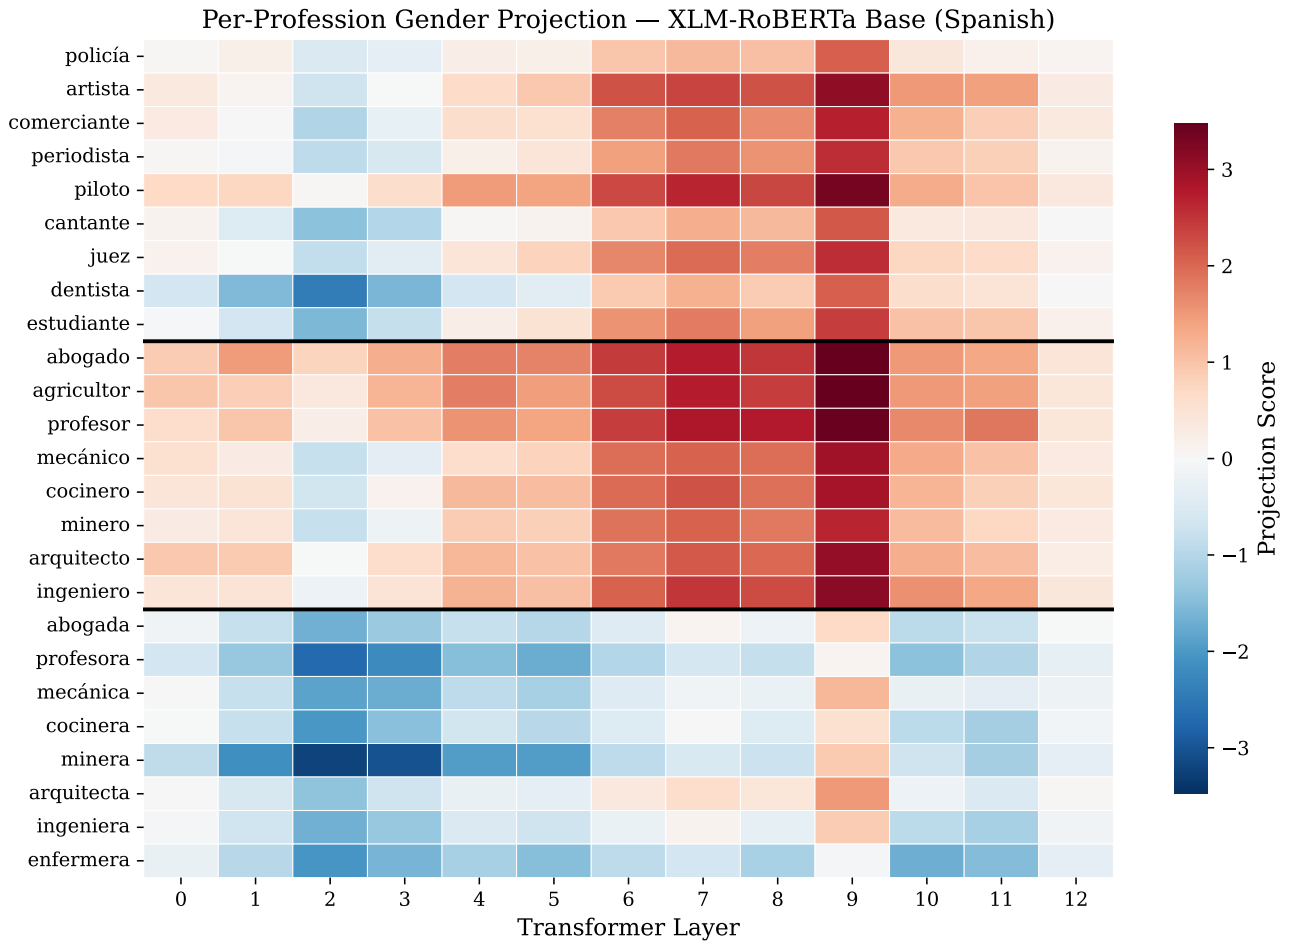

Figure 5: Full per profession heatmap for XLM-RoBERTa Base. Horizontal lines separate Group A (invariant), Group B masculine, and Group B feminine professions.

Table 4: Cross lingual summary: layer wise projection statistics for 12-layer models.

Lang.	**Model**	**Peak Layer**	**Peak Proj.**	**Min ρ**	n (prof.)
Spanish	XLM-R Base	9	+2.15	0.991	25
Arabic	XLM-R Base	9	+2.64	0.964	23
Tigrinya	XLM-R Base	9	+2.99	0.992	30

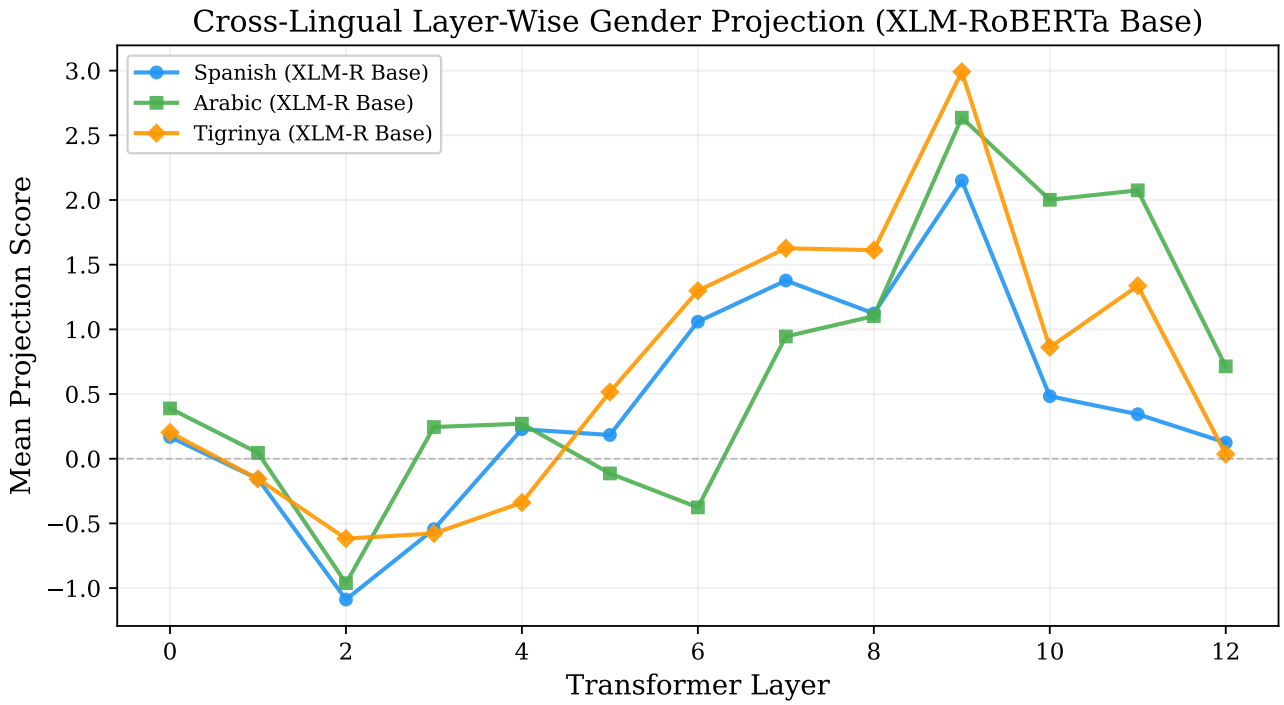

Figure 6: Cross lingual layer wise mean projection for XLM-RoBERTa Base across Spanish, Arabic, and Tigrinya. Remarkably, the architecture encodes bias with an identical trajectory (peaking exactly at Layer 9) regardless of the language's resource level.

5 Discussion and Conclusions

5.1 Grammatical Neutrality vs. Societal Frequency

Multilingual transformers learn morphological neutrality without explicit supervision, locating invariant professions (*dentista*, *cantante*) near the geometric origin. However, statistical frequency in the pretraining corpus overrides grammatical rules for historically male dominated roles. For example, XLM-R Base projects the invariant *comerciante* (+2.69) similarly to the explicitly masculine *mecánico* (+2.94). Only mDeBERTa-v3 suppresses this effect (−0.03), indicating that architectural design also has an influence in the conflict between grammar and frequency.

5.2 The Masculine Default

The asymmetry analysis shows that masculine forms occupy a dominant position as the default representation of a profession. Feminine forms correctly orient toward the female centroid but remain closer to the origin. By preserving this data asymmetry or imbalance, the models treat masculine professions as more semantically central, carrying practical implications for certain applications like search retrieval or automated text analysis, and the interpretability research.

5.3 Architecture over Pretraining Data

Our results emphasize architecture as a dominant factor in bias encoding. mDeBERTa-v3 reverses the bias direction relative to RoBERTa family models, preserving relative profession rankings ($\rho > 0.97$) but flipping their geometric orientation. We attribute this to Disentangled Attention, by separating content from positional context, the model prevents societal frequencies from having too much influence on the gender direction.

Similarly, the cross lingual analysis shows that semantic structures depend less on pretraining volume than often assumed. XLM-R Base encodes bias with an identical structural trajectory (peaking at Layer 9) across Spanish, Arabic, and Tigrinya, despite massive disparities in corpus size, opening the question wether the attention mechanism functions as a universal template for how abstractions like gender are structured.

5.4 Final Remarks and Limitations

Since we only looked at base embeddings without fine tuning, it's possible these bias patterns could change once the models are adapted for specific tasks. Moreover, our anchor sets and Spanish context corpus may not capture all domains of language use.

In conclusion, multilingual transformers encode measurable geometric gender bias that peaks in the middle semantic layers. Evaluating fairness exclusively at the final layer might underestimate magnitude. Our results suggest that measuring geometric structures across layers is more reliable. Furthermore, architectural choices like attention factorisation seem to mitigate social bias, bypassing the need for explicit debiasing pipelines.

References

Bolukbasi, Tolga et al. (2016). "Man is to computer programmer as woman is to homemaker? Debiasing word embeddings". In: *Advances in Neural Information Processing Systems*. Vol. 29.

Bommasani, Rishi, Kelly Davis, and Claire Cardie (2020). "Interpreting pretrained contextualized representations via reductions to static embeddings". In: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*. Association for Computational Linguistics, pp. 4758–4781.

Caliskan, Aylin, Joanna J. Bryson, and Arvind Narayanan (2017). "Semantics derived automatically from language corpora contain human-like biases". In: *Science* 356.6334, pp. 183–186.

Conneau, Alexis et al. (2020). "Unsupervised cross-lingual representation learning at scale". In: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*. Association for Computational Linguistics, pp. 8440–8451.

Devlin, Jacob et al. (2019). "BERT: Pre-training of deep bidirectional transformers for language understanding". In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics*. Association for Computational Linguistics, pp. 4171–4186.

Garg, Nikhil et al. (2018). "Word embeddings quantify 100 years of gender and ethnic stereotypes". In: *Proceedings of the National Academy of Sciences* 115.16, E3635–E3644.

Gonen, Hila and Yoav Goldberg (2019). "Lipstick on a pig: Debiasing methods cover up systematic gender biases in word embeddings but do not remove them". In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Association for Computational Linguistics.

Guo, Wei and Aylin Caliskan (2021). "Detecting emergent intersectional biases: Contextualized word embeddings contain a distribution of human-like biases". In: *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*, pp. 122–133.

He, Pengcheng, Jianfeng Gao, and Weizhu Chen (2023). "DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing". In: *International Conference on Learning Representations*.

He, Pengcheng et al. (2021). "DeBERTa: Decoding-enhanced BERT with disentangled attention". In: *International Conference on Learning Representations*.

Lauscher, Anne et al. (2020). "From zero to hero: On the limitations of zero-shot language transfer with multilingual Transformers". In: *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, pp. 4483–4499. DOI: 10.18653/v1/2020.emnlp-main.363.

Liang, Davis et al. (2023). "XLM-V: Overcoming the vocabulary bottleneck in multilingual masked language models". In: *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, pp. 13142–13152.

Rogers, Anna, Olga Kovaleva, and Anna Rumshisky (2020). "A primer in BERTology: What we know about how BERT works". In: *Transactions of the Association for Computational Linguistics* 8, pp. 842–866.

Spearman, Charles (1904). "The proof and measurement of association between two things". In: *The American Journal of Psychology* 15.1, pp. 72–101.

Zhao, Jieyu et al. (2019). "Gender bias in contextualized word embeddings". In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics*, pp. 629–634.

Zhao, Jieyu et al. (2020). "Gender bias in multilingual embeddings and cross-lingual transfer". In: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*. Association for Computational Linguistics, pp. 2896–2907. DOI: 10.18653/v1/2020.acl-main.260.

Appendix A. Reproducibility Statement

To facilitate full reproducibility, the source code, data parsing scripts, and complete evaluation pipeline used in this study are publicly available in the following GitHub repository: https://github.com/JuanArango99/ml4nlu-gender-bias.

Appendix B. Statement on the Use of Artificial Intelligence

In accordance with the academic guidelines for the module *Machine Learning for Natural Language Understanding*, this section discloses the use of generative Artificial Intelligence (AI) in the preparation of this term paper.

The main author (Juan Arango) utilized Google Gemini for the following tasks:

- Assisting in the implementation of the data loading pipeline (`data_loader.py`), the geometric bias calculation framework, and part of the code for the visualization generation script.
- Textual Refinement, proofreading and improving the formal academic phrasing of the Abstract, Introduction, and Methodology sections to ensure clarity and LaTeX formatting standards.

All AI generated suggestions, code snippets, and textual edits were critically reviewed, verified, and modified by the main author. The final responsibility for the technical accuracy and academic integrity of this work rests solely with the main author.

Appendix C. Notable Individual Model Trajectories

While the main text aggregates layer wise projections or restricts analysis to XLM-RoBERTa Base to uncover structural consistency, individual models show distinct behaviors. We use four examples below.

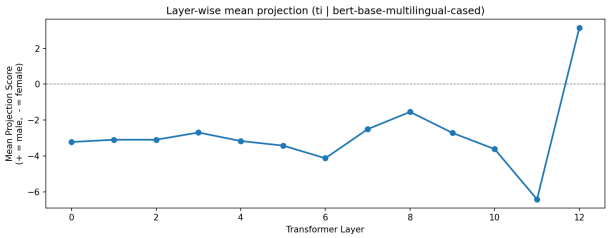

Figure 7: Tigrinya (mBERT): Extreme low-resource collapse. With insufficient pretraining data, mBERT fails to generate a societal frequency signal, collapsing into arbitrary negative projections.

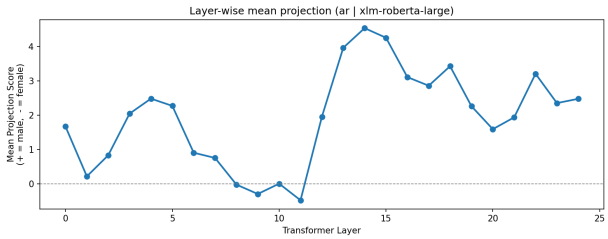

Figure 8: Arabic (XLM-R Large): Late stage peaking. While Base models peak near layer 9, this Large model peaks at layer 14 with a +4.54 magnitude.

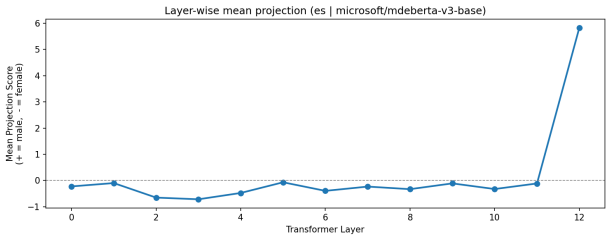

Figure 9: Spanish (mDeBERTa-v3): Disentangled Attention reversal. The projection score stays negative across all middle layers, reversing the RoBERTa upward curve.

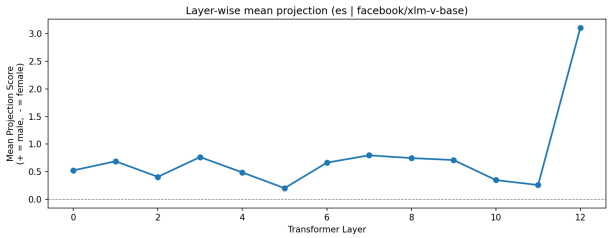

Figure 10: Spanish (XLM-V Base): Final layer artifact spike. The bias score exhibits an unnatural spike at layer 12 (+3.11) possibly due to token level output preparation.

Appendix D. Evaluated Vocabulary (Spanish)

This section details the explicit Spanish word sets used to construct the gender subspace and evaluate profession embeddings. The anchor sets (Table 5) capture fundamental gender semantics, while the profession sets (Table 6) are divided into morphologically invariant (Group A) and gender inflected (Group B) terms.

Table 5: Spanish gender anchor sets used to define the semantic subspace.

Masculine Anchors	**Feminine Anchors**
hombre *(man)*	mujer *(woman)*
niño *(boy)*	niña *(girl)*
padre *(father)*	madre *(mother)*
él *(he)*	ella *(she)*
hermano *(brother)*	hermana *(sister)*
esposo *(husband)*	esposa *(wife)*
abuelo *(grandfather)*	abuela *(grandmother)*
hijo *(son)*	hija *(daughter)*
tío *(uncle)*	tía *(aunt)*
varón *(male)*	señora *(lady)*
chico *(boy/guy)*	chica *(girl/gal)*

Table 6: Spanish profession terms evaluated in this study, categorized by morphological group.

Group A (Invariant)	**Group B (Masculine)**	**Group B (Feminine)**
policía *(police)*	abogado *(lawyer)*	abogada *(lawyer)*
recepcionista *(receptionist)*	enfermero *(nurse)*	enfermera *(nurse)*
artista *(artist)*	agricultor *(farmer)*	agricultora *(farmer)*
comerciante *(merchant)*	profesor *(professor)*	profesora *(professor)*
periodista *(journalist)*	mecánico *(mechanic)*	mecánica *(mechanic)*
piloto *(pilot)*	cocinero *(cook)*	cocinera *(cook)*
cantante *(singer)*	limpiador *(cleaner)*	limpiadora *(cleaner)*
juez *(judge)*	minero *(miner)*	minera *(miner)*
dentista *(dentist)*	arquitecto *(architect)*	arquitecta *(architect)*
estudiante *(student)*	ingeniero *(engineer)*	ingeniera *(engineer)*